

BACHELOR THESIS

# Roadster High Availability

*Patrik Wenger, Manuel Schuler*

for industry client  
mindclue GmbH

supervised by  
Prof. Farhad Mehta

Fall semester 2016

## **Abstract**

TODO introduction

TODO approach and technologies

TODO result

# Declaration of Originality

We hereby confirm that we are the sole authors of this document and the described changes to the Roadster framework and libraries developed.

TODO any usage agreements or license

# Acknowledgements

TODO anyone we'd like to thank

# Management Summary

# Context

## Initial Situation

Roadster is mindclue GmbH's in-house framework to build modern monitoring and controlling applications in different fields such as traffic systems, energy, and water supply. It is written in Ruby, a modern and expressive scripting language, and is built on a shared-nothing architecture to avoid a whole class of concurrency and scalability issues found in traditional application architectures.

Although considered to be the next generation of its kind, it still lacks important features such as the ability to be run as a cluster on multiple nodes, high availability, and secure network communications.

## Goals

Adding the aforementioned, missing features to form the next version of the framework would mean a distinct advantage for mindclue GmbH and thus increase its competitiveness in its sector.

Planning the exact architectural changes and additions, as well as performing the implementation is the students' goal for this bachelor thesis. Using engineering methodology practiced at HSR, solutions for particular problems will be worked out and the best fitting one will be chosen.

Although not exactly part of the requirements, spreading knowledge about Roadster's architecture and code basis is also in the interest of the client, as Andy Rohr is currently the framework's only developer and thus a single point of failure in an increasingly important piece of software.

## Software Development Process

TODO take RUP into glossary

RUP is used to plan and manage this term project. It's an iterative, structured, yet flexible development process which suits this kind of project. At HSR, it's taught as part of the Software Engineering courses and is thus a primary candidate.

Another candidate was Scrum, which we decided against as it's only feasible with teams of three to nine developers.

## Project Management Infrastructure

The source code of this document and all of our code contributions are hosted on GitHub. The students will organize and perform their work directly on the site as far as possible. This means creating a Project board for each of the development phases, creating, assigning, and closing issues, as well as using the Wiki feature to plan and document meetings with the professor and the client.

Time tracking, as required by the process for bachelor theses [**hsr:thesis-rules**], are done externally on Everhour.

## Personal Goals

Next to excelling in this bachelor thesis, the students also hope to create a reusable open-source library as a by-product. The intention is that the library makes certain ØMQ-based communication protocols readily available for other developers facing the same problems.

## Project Phases

TODO describe this phase in retrospection

### Inception

TODO include Gantt chart for this phase  
TODO describe this phase in retrospection

### Elaboration

TODO include Gantt chart for this phase  
TODO describe this phase in retrospection

### Construction

TODO include Gantt chart for this phase  
TODO describe this phase in retrospection

### Transition

TODO include Gantt chart for this phase  
TODO describe this phase in retrospection

## Results

TODO describe results

# Contents

<b>I</b>	<b>Technical Report</b>	<b>1</b>
<b>1</b>	<b>Scope</b>	<b>2</b>
	Motivation . . . . .	2
	Open-Source Engagement . . . . .	2
	Initial Situation . . . . .	3
	mindclue GmbH . . . . .	3
	ØMQ . . . . .	3
	Software Architecture . . . . .	4
	Goals . . . . .	5
<b>2</b>	<b>Requirements</b>	<b>7</b>
	Priorities . . . . .	7
	Functional . . . . .	7
	Cluster . . . . .	7
	Single Level HA . . . . .	8
	Multi Level HA . . . . .	8
	Persistence Synchronization . . . . .	8
	Security . . . . .	8
	OPC UA HA . . . . .	9
	Use Cases . . . . .	9
	Non-Functional Requirements . . . . .	9
	Simplicity . . . . .	9
	Testing . . . . .	9
	Coding Guidelines . . . . .	9
<b>3</b>	<b>Methodology</b>	<b>10</b>
	Port to new ØMQ library . . . . .	10
	Cluster . . . . .	10
	Aspects . . . . .	10
	DIM Synchronization . . . . .	11
	Node Typology Definition . . . . .	13
	Message Routing . . . . .	14
	High Availability . . . . .	15
	Single Level . . . . .	15
	Multi Level . . . . .	15
	Persistence Synchronization . . . . .	16
	Aspects . . . . .	16
	Variants . . . . .	16
	Chosen Variant . . . . .	18
	Security . . . . .	18
	OPC UA Interface: High Availability . . . . .	18
<b>4</b>	<b>Results</b>	<b>19</b>
	Port . . . . .	19



Cluster . . . . .	19
High Availability . . . . .	19
Persistence Synchronization . . . . .	20
Security . . . . .	20
OPC UA Interface: High Availability . . . . .	20
<b>5 Discussion</b>	<b>21</b>
Value Added . . . . .	21
Limitations . . . . .	21
Business Benefits . . . . .	21
Ideas for Improvement . . . . .	21
<b>6 Conclusion</b>	<b>23</b>
 <b>II Appendix</b>	 <b>25</b>
<b>A Self Reflection</b>	<b>26</b>
<b>B Task Description</b>	<b>27</b>
<b>C License</b>	<b>28</b>
<b>D Project Plan</b>	<b>29</b>
Organization . . . . .	29
<b>E ZMQ</b>	<b>30</b>
<b>F Infrastructural Problems</b>	<b>31</b>
Project Management Software . . . . .	31

# List of Figures

1.1	Roadster's software architecture . . . . .	5
1.2	Roadster's communication layers . . . . .	6

# List of Tables

# Listings

## Part I

# Technical Report

# Chapter 1

## Scope

The technical goals of this bachelor thesis include extending mindclue GmbH's Roadster framework by adding features such as clustering, high availability and transport security. This chapter outlines the general scope of this project.

## Motivation

TODO Why do we care about this thesis? Why are we interested?

We're thrilled about the following technologies and software design patterns:

- ØMQ
- Ruby
- Actor Model
- High Availability

Coming from different backgrounds and having different degrees of experience in each of the above technologies, we can't wait to learn more about them and put them to actual use. The fact that the product of this bachelor thesis is most likely going to be used in the real world only adds to the excitement.

In addition to that, we look at this bachelor thesis as an opportunity to become more fluent in English, as well as a way to improve our skills in crafting scientific documents using L<sup>A</sup>T<sub>E</sub>X.

Depending on how we perform together as a team, further collaboration might result in the future, either between the students themselves, or between the students and the client. Even if not, this project will serve as a valuable reference for future job hunting.

Last but not least, we feel like Prof. Dr. Mehta is a respected and competent teacher whose opinions we highly value. Due to his polite parlance, discussing project matters, both of the management and the technical kind, has always been an enrichment.

## Open-Source Engagement

TODO CZTop, cztop-patterns (as mentioned in the Mgmt Summary)

# Initial Situation

## mindclue GmbH

### About

Die Firma mindclue GmbH mit Sitz in Ziegelbrücke GL/SG stellt für ihren Partner REMTEC AG für Steuerungen, Regelungen und Überwachungen von technischen Prozessen, komplette SCADA- und Steuerungssysteme her. Dabei sind sie in den Bereichen Betriebs- und Sicherheitsausrüstung für Nationalstrassen, Trinkwasser- versorgungen, im Energiesektor und vielen Spezialbereichen tätig. Dabei setzen sie auf ihre eigens entwickelte Applikation - Roadster.

### Roadster

Roadster ist eine Ereignis gesteuerte Applikation geschrieben in Ruby. Sie ist die eigentliche SCADA Applikation und wird bereits in einer Vielzahl von Tunnelanlagen in der Schweiz eingesetzt und dient zur Überwachung der einzelnen Komponenten im Tunnel. Durch den modularen Aufbau kann jede Roadster Applikation sich selbst laufen - Autonom. Sprich jeder Roadster hat seinen eigenen Webserver und Datenverwaltung.

### AS - UeLS

Roadster ist einer von vielen AS Knoten eines UeLS. Die Kommunikation zwischen AS und AR geschieht über "OPC UA".

### ØMQ

*For a more detailed introduction, see Appendix E.* To understand Roadster's architecture and the rest of this document, it's helpful to understand the basics of ØMQ (sometimes written as ZeroMQ or simply ZMQ) first. This is a brief introduction to ØMQ for the unfamiliar reader.

ØMQ is a MOM implemented as an open source library, that is, it doesn't require a dedicated broker. Instead, it offers sockets with an abstract interface similar to BSD sockets. Different types of sockets are used for different messaging patterns such as request-reply, publish-subscribe, and push-pull.

A single socket can bind/connect to multiple endpoints, which allows ØMQ to use round-robin on the sender side, and fair-queueing on the receiver side, where applicable. It doesn't matter whether the communication happens in-process (between threads), inter-process (e.g. over Unix Domain Sockets), or inter-node (e.g. over TCP/PGM/TIPC), since the transport is completely abstracted away. The same goes for connection handling; an arbitrary amount of connections is handled over a single socket and reconnecting after short network failures is done transparently.

ØMQ is lightweight and provides extremely low latencies, which means it can also be used as the fabric of concurrent applications, e.g. for the actor model. In case of the TCP transport, it incorporates advanced techniques such as smart message batching to achieve significantly higher throughputs than with raw TCP or other MOM solutions [2, Figure 2, Middleware evaluation and prototyping, p. 4].

To build a solution with ØMQ, its sockets are used as building blocks to design custom message flows. Certain patterns are used to achieve reliability with respect to the failure types that need

to be addressed in particular. The `zguide`<sup>1</sup> explains best practices, including commonly needed, resilient messaging patterns.

The above characteristics make ØMQ a valuable asset when it comes to building robust, distributed high-performance systems.

## Transport Security

Since version 4.0, ØMQ boasts state of the art encryption and authentication, based on the excellent and highly renown NaCl<sup>2</sup> library.

## Data Serialization

Data serialization is outside the scope of ØMQ. To fill the gap, one typically uses another library such as `MsgPack`<sup>3</sup>, `Protocol Buffers`<sup>4</sup>, or even a programming language's built-in object serialization support<sup>5</sup>.

## CZMQ

CZMQ is a high-level abstraction layer for ØMQ. It makes working with the ØMQ library more expressive and allows for better portability. It also provides additional functionality such as a reactor, a simple actor implementation, as well as utilities for certificate and authentication handling, and LAN node discovery. This is the recommended way of using ØMQ nowadays.

## Software Architecture

TODO more

Roadster is event-driven and built on the Actor model, meaning it exhibits a shared-nothing architecture. Each Roadster node runs a number of Ruby processes which communicate via ØMQ sockets. The key here is communication:

“Don’t communicate by sharing state; share state by communicating.”

Running multiple, loosely coupled processes (actors) allows leveraging the full potential of modern multi-core processors, while avoiding a whole class of traditional concurrency problems.

Every Roadster node runs a group of actors:

**CORE:** It is responsible to start the other actors. It also plays a key role in keeping state in all actors synchronized, being the source of truth.

**COMM:** A bunch of COMM actors communicate with the outside world of a node. This can be different kind of PLCs or higher level monitoring systems. Each COMM actor uses an adapter specifically written for a single communication protocol.

**STORAGE:** This actor is used when information needs to be persisted, such as time series or event journals. It’s the interface to a key-value store.

**LOGGER:** This actor collects logging data and sends it to whatever target is configured, be it STDOUT, a file, or a syslog server.

---

<sup>1</sup><http://zguide.zeromq.org/>

<sup>2</sup><http://nacl.cr.yp.to>

<sup>3</sup><http://msgpack.org>

<sup>4</sup><https://developers.google.com/protocol-buffers/>

<sup>5</sup>such as Ruby’s marshalling support: <http://ruby-doc.org/core/Marshal.html>



Figure 1.1 illustrates Roadster's architecture.

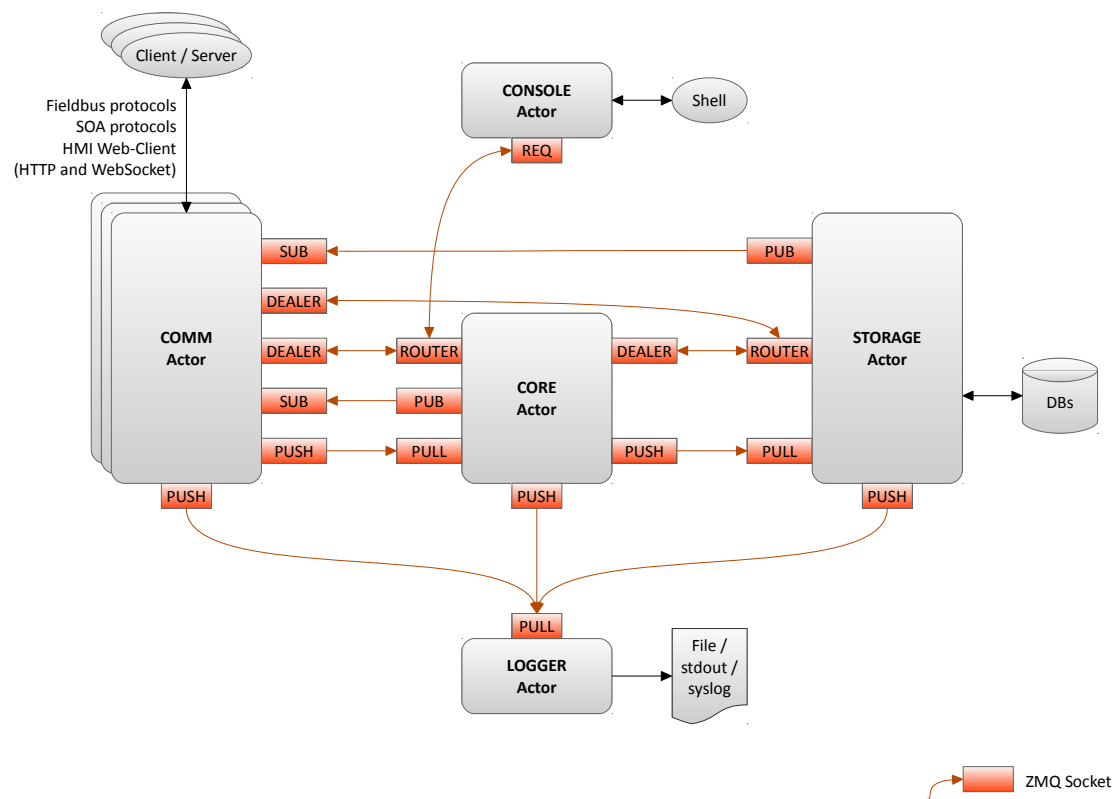


Figure 1.1: Roadster's software architecture

## Communication Layers

TODO briefly explain layers. Figure 1.2 illustrates the three layers.

## Goals

TODO mandatory goals

## Optional Goals

TODO optional goals

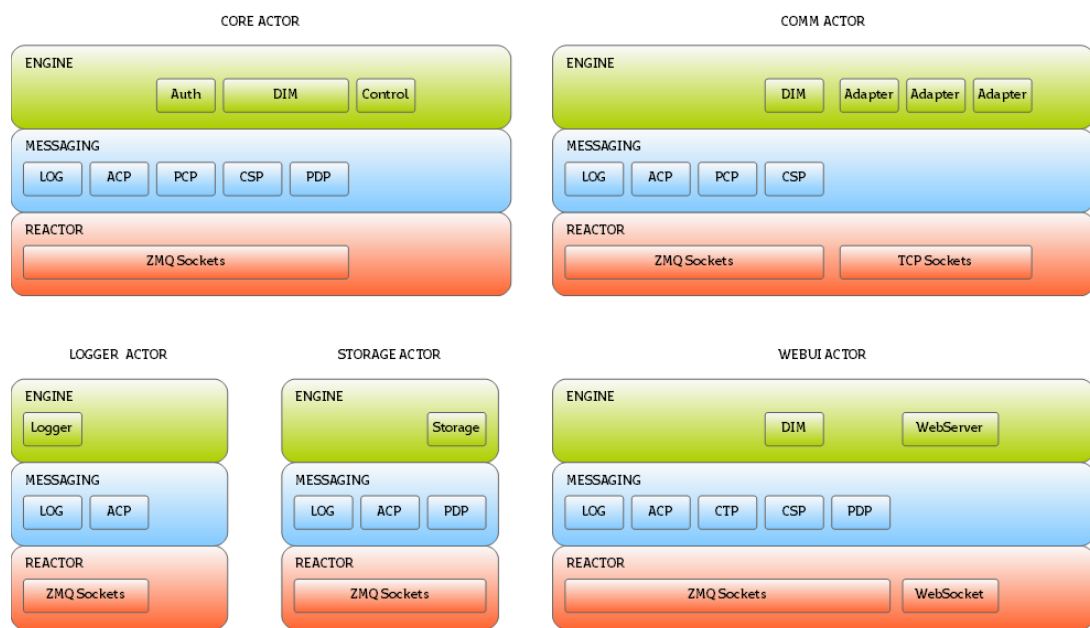


Figure 1.2: Roadster's communication layers

TODO use PDF

## Chapter 2

# Requirements

TODO the requirements

## Priorities

In descending priority:

1. multi-node CSP
2. single-level HA
3. multi-level HA
4. persistence synchronization
5. security
6. OPC UA HA (optional)

The following sections explain the requirements in greater detail.

## Functional

### Cluster

This could also be called "Multi-node CSP".

- this is to allow running Roadster in a hierarchical setup
- new COMM actors for inter node communication
- usually 2 (or 3) levels of Roadster nodes
- common cases:
  - - single level, single node (legacy)
  - - single level HA
  - - multi level, HA at root node only
- exotic cases:
  - - multi level, HA at bottom
  - - multi level, HA in middle

- every subtree can live on autonomously
- only node A has write access to values on A (to avoid uncertain situations involving race conditions), e.g.:
  - - a forced value coming from the web UI comes through a command,
  - - routed to the relevant node, where it is applied,
  - - and then synced (up via DEALER and down via PUB, we suppose)



## Single Level HA

This is where there's a node pair directly connected to a PLC. Both nodes have read/write access to the PLC, but only one of the nodes (the active one) must do so. The nodes must automatically find consensus on who's active. The passive one must automatically take over in case the active one is confirmed to be dead.

TODO the kinds of failures we want to be able to handle: exactly hardware/software failure of the primary node, and network failure (stated by the Task Description)

## Multi Level HA

This is where a node pair is the parent of one or more other nodes (subnodes).

TODO the kinds of failures we want to be able to handle: exactly hardware/software failure of the primary node, and network failure (stated by the Task Description)

## Persistence Synchronization

This is about the synchronization of the TokyoCabinet databases. Data flow is from south to north (towards the root node), so the root node collects and maintains a replication of the persisted data of all subnodes, recursively.

- autonomous
- not same as CHP
- 100% consistency is not important
- data only flows from bottom to top
- TC keys contain timestamp

## Security

- transport needs to be secure (encrypted and authenticated)
- TODO verify requirements with Andy (we didn't really discuss this during the meeting)

- this requirement comes as the last mandatory goal not because it's insignificant, but because it's easy to enable transport level security on ZMQ sockets, and it would just interfere with the previous development

## OPC UA HA

- provide standardized interface upwards from HA pair

## Use Cases

TODO maybe there are any?

## Non-Functional Requirements

### Simplicity

The two reoccurring patterns that surfaced during the requirements gathering meeting were:

1. KISS principle – “Keep it simple, stupid.” Simplicity is favored, unnecessary complexity should be avoided.
2. No premature optimization since it's the root of all evil.<sup>1</sup>

### Testing

- we write unit tests for our own contributions
- we test the integrated result in a close-to-reality setup

## Coding Guidelines

The coding guidelines desired by the client are basically the ones written down in the popular Ruby style guide [1], with the following differences or special remarks:

- method calls: only use parenthesis when needed, even with arguments (as opposed to <sup>2</sup>)
- 2 blank lines before method definition (slightly extending <sup>3</sup>)
- YARD API doc, 1 blank comment line before param documentation, one blank comment line before code (ignoring <sup>4</sup>)
- Ruby 1.9 symbol keys are wanted (e.g. `foo: "bar"`, `baz: 42` instead of `:foo => "bar"`, `:baz => 42`, just like <sup>5</sup>)
- align multiple assignments so there's a column of equal signs

<sup>1</sup>Quote by Donald Knuth: “Premature optimization is the root of all evil.”

<sup>2</sup><https://github.com/bbatsov/ruby-style-guide#method-invocation-parens>

<sup>3</sup><https://github.com/bbatsov/ruby-style-guide#empty-lines-between-methods>

<sup>4</sup><https://github.com/bbatsov/ruby-style-guide#rdoc-conventions>

<sup>5</sup><https://github.com/bbatsov/ruby-style-guide#hash-literals>

## Chapter 3

# Methodology

TODO what have we done to arrive at the goal (should be reproducible)  
TODO this is probably what we know as "Concept"

### Port to new ØMQ library

TODO justify why port is needed right at the beginning (exclude faults from unmaintained ffi-rmq gem, encryption is needed anyway, all the other tasks involve communication over ZMQ)  
TODO explain binding options out there, why CZTop (including difference between ZMQ and CZMQ)  
TODO explain preliminary task of adding support for the ZMQ options ZMQ\_FD and ZMQ\_EVENTS in CZTop  
TODO explain concept of exchanging ffi-rmq with CZTop

### Cluster

TODO describe scribble (chp.pdf)

### Aspects

Adding cluster functionality to Roadster involves the following aspects:

- node topology DSL  
This DSL also has to provide means to define the roles/functionality of each node, e.g. the set of COMM actors running on a particular node
- DIM synchronization
- message routing
- What needs to be done a WebUI user wants to e.g. change some value on a PLC, possibly on a remote node? Is it completely handled via DIM or do we need message routing?

## DIM Synchronization

TODO election/design of appropriate protocol

### The Existing CSP in a Nutshell

*This is a brief introduction/refreshers for the Clone State Pattern implemented by Roadster. Although Roadster actually sends serialized instances of CSP message classes to fulfill this protocol, for better readability the Zguide's canonical nomenclature of Clone Pattern messages will be used.*

The existing CSP is closely related to the Clone Pattern from the Zguide<sup>1</sup>. Its goal is to keep a state (a list of key-value pairs) in sync across a set of participants. To greatly reduce the complexity, it's not decentralized: There's a server part which serves as the single source of truth.

The server uses a ROUTER, a PULL, and a PUB socket; each client a DEALER, a PUSH, and a SUB socket. The protocol consists of three distinct messages flows:

**Snapshots:** Requesting and receiving the complete, current snapshot of the state (all key-value pairs). This happens via a ROUTER/DEALER pair of sockets. The request message consists solely of the humorously named ICANHAZ command. The response is the complete set of KVSET messages so a late-joining (or previously disconnected) client can rebuild the current snapshot.

**Upstream updates:** Updates always originate from clients and are sent to the server via a PUSH/PULL pair of sockets. These are KVSET messages.

**Downstream updates:** After being applied to the server's copy of the state, updates get a sequence number and are published back to all clients. This happens via the PUB socket and uses KVPUB messages.

By making all updates go through the server, a total order is enforced, which is crucial to keep the state consistent across all clients.

To avoid risking a gap between requesting the current snapshot and subscribing to updates, a client actually subscribes to the updates first, then gets the snapshot, and then starts reading the updates from the socket (which has been queueing updates in the meantime, if any). Updates that are older or the same age as the received snapshot are skipped, and only successive updates are applied (tested by comparing the sequence numbers).

Because message loss via the third message flow (PUB-SUB) is unlikely but theoretically possible, the client checks for gaps in the sequence number of each KVPUB message. If a gap is detected, the current state is discarded and a complete resynchronization happens. This is brutal, but is very simple and thus robust; there's no complexity that would leave room for nasty corner cases.

Keys can be treated hierarchically (e.g. `topic.subtopic.key`) and thus, a client can optionally subscribe to only a particular subtree. This is useful when the number of client grows and not all of the state needs to be on every client. In that case, the topic of interest is sent by the client along with the ICANHAZ message.

### What's Missing

TODO it has to work across several nodes

TODO it has to be able to handle HA supernodes

---

<sup>1</sup>FIXME

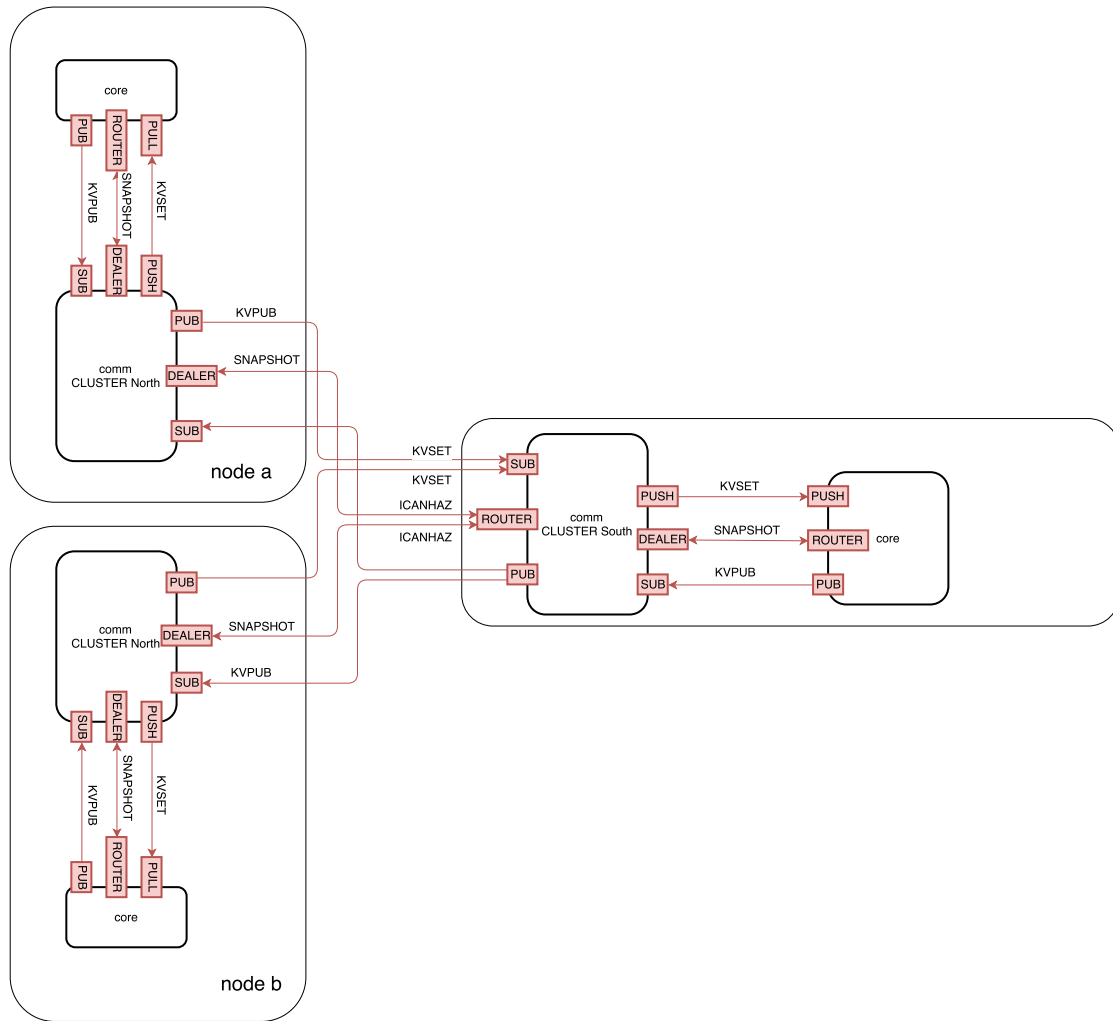
TODO explain Clustered Hashmap Protocol (?)

- PCP: use DIM to know node tree and determine next hop for (dialog or fire+forget) messages
- decide on sync variant
  - variant 1
    - \* always sync on self-subtree only
    - \* con: no copy of remaining tree
  - variant 2:
    - \* always sync on complete tree
    - \* get snapshot and merge own subtree
    - \* (this should probably be the first step)
  - variant 3:
    - \* make it configurable: either sync on subtree or complete tree
    - \* (this should probably be the second step, if at all)

#### **What's needed**

- we need two new COMM actors
- they sync between super node and sub node
- they do something closely related to the existing CSP
- future oriented: because of the HA requirement, ideas from the CHP are integrated, such as using PUB-SUB (instead of PUSH-PULL) for inter-node KVSET messages, so both super nodes (in HA setup) hear updates
- for intra-node KVSET messages, PUSH-PULL is OK and can be left unchanged





## Node Typology Definition

- node topology in DSL, static file (e.g. topology.conf.rb) shared on all nodes, read by each actor on startup
- specific config file on each node (conf.rb) knows its own place in topology (through `conf.system_id`)
- maybe a HA pair is one DIM object, has one name, but two IP addresses (primary and backup, in order)

```
# * basic method to add a node: #add_node(ID, south_facing_bind_endpoint)
# * it takes a block for defining subnodes

#####
# without HA:

conf.nodes do |map|
  map.add_node("root", "tcp://10.0.0.1:5000") do |map|
    map.add_node("subnode_a", "tcp://10.0.0.10:5000")
    map.add_node("subnode_b", "tcp://10.0.0.11:5000")
  end
end

# subnode_a can infer its endpoints from its position in the tree:
conf.system_id = "nodes.root.subnode_a"
```

```

#=> this node is "subnode_a"
#=> its IP address is 10.0.0.10
#=> north facing COMM actor's bind port is 5001
#=> south facing COMM actor's bind port is 5000
#=> north facing COMM actor will connect to "root" node on "tcp://10.0.0.1:5000"

#####
# later with HA:

conf.nodes do |map|
  map.add_ha_pair("root", "tcp://10.0.0.1:5000", "tcp://10.0.0.2:5000") do |map|
    map.add_node("subnode_a", "tcp://10.0.0.10:5000")
    map.add_node("subnode_b", "tcp://10.0.0.11:5000")
  end
end

# subnodeA can infer its endpoints from its position in the tree:
conf.system_id = "nodes.root.subnode_a"
#=> this node is "subnode_a"
#=> its IP address is 10.0.0.10
#=> north facing COMM actor's bind port is 5001
#=> south facing COMM actor's bind port is 5000
#=> north facing COMM actor will connect to "root" HA pair on "tcp←
    ↳ ://10.0.0.1:5000" OR "tcp://10.0.0.2:5000" (Lazy Pirate algorithm)

# for primary root:
conf.system_id = "nodes.root[primary]"

#####
# within ba-roadster-app's lib/domain/domain.rb file:
#
# Idea for node topology definition and assigning roles (features/adapters) to
# different kinds of nodes.

module Roadster
  module Domain::Model

    build do
      nodes do
        node "root" do # or maybe ha_node or bstar_node
          endpoint "tcp://10.0.0.1:5000", "tcp://10.0.0.2:5000"
          label 'BA Roadster App'
          desc 'Sample application for experimenting and developing the new ←
            ↳ features within the scope of the Bachelor Thesis of Patrik ←
            ↳ Wenger and Manuel Schuler at HSR.'

          load_conf ::Conf::AccessControl
          load_conf ::Conf::Objects
          load_conf ::Conf::Navigation

          node "subnode_a" do
            endpoint "tcp://10.0.0.1:5000"
            load_conf ::Conf::Adapters
            # load_conf ...
          end
        end
      end
    end

  end # Domain::Model
end # Roadster

```

## Message Routing

TODO message routing (end-to-end routing with identity/identities as prepended message frame?, should be simpler and more efficient than hop-by-hop routing)

# High Availability

TODO we have two different kinds of HA

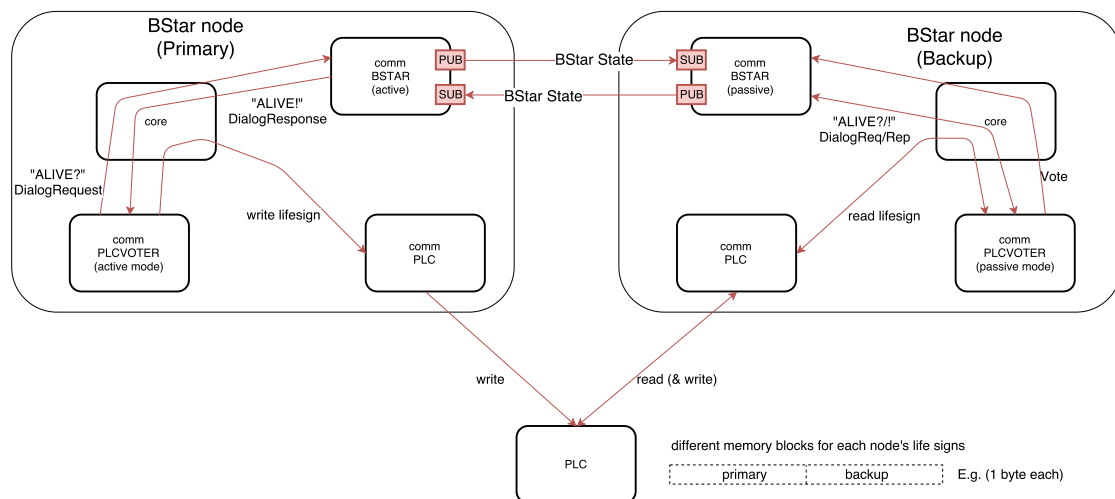
TODO explain how the failures we're required to be able to handle can be handled

TODO explain similarities between the two kinds of HA

TODO case generation in case of failover or passive down

## Single Level

- this is different from what's described in the zguide because the concept of client requests is missing here (PLCs don't request anything)
- life sign from one node to the other through some continually updated PLC value
- don't send life signs when, for example, CORE is dead!
- mark active HA peer in DIM

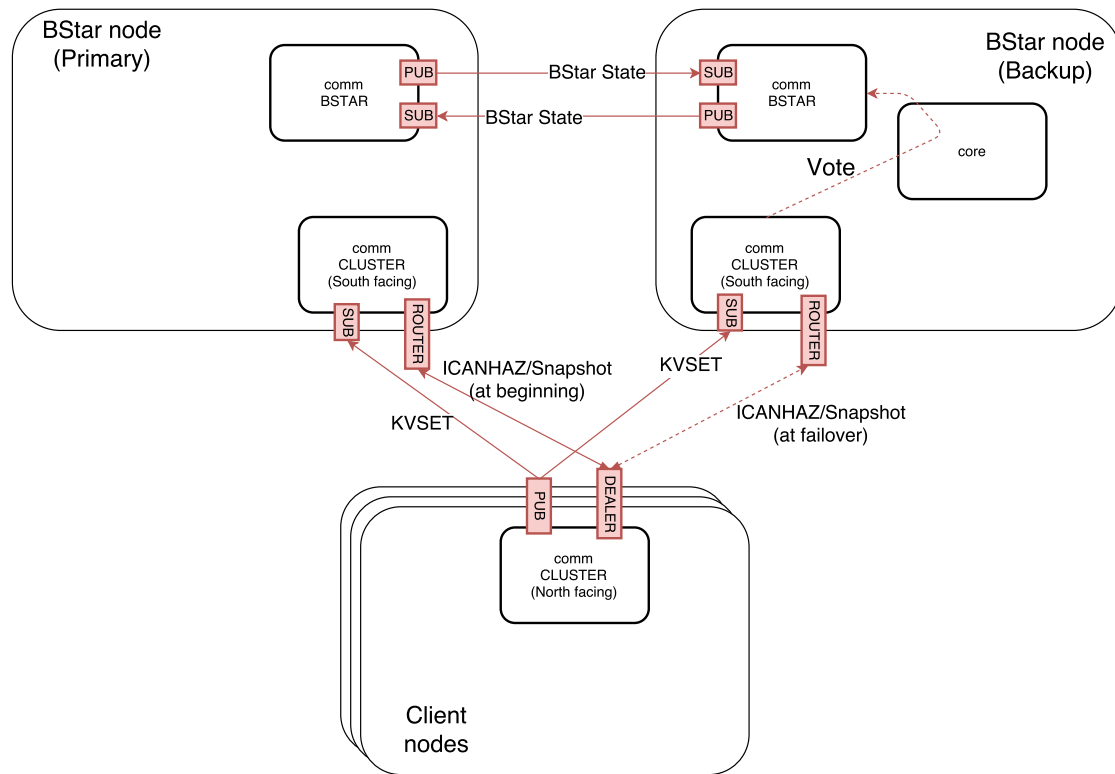


## Multi Level

TODO explain why is this one different from SL-HA

TODO Finding consensus should be easier here, as it's closely related to the CHP described in the zguide.

TODO integrate findings from scribble (ML-HA.pdf)



## Persistence Synchronization

- super node requests for delta of TC periodically

### Aspects

There are multiple aspects involved in persistence synchronization:

**Delta:** How does one get the initial delta of updates since last synchronization?

**Updates:** Further updates, one-by-one. This is only needed in case the solution aims for real-time synchronization.

**HA peer sync:** How does the inactive HA peer get updated? Of course, this only matters when the supernode is HA pair.

### Variants

There are multiple variants to achieve the needed functionality.

#### Polling only

The supernode just periodically request persistence deltas. This would be handled over a DEALER/ROUTER pair of sockets. The nice thing about this variant is that the subnode only has to do one thing, which is responding to requests from the supernode(s); it doesn't have to proactively send any updates after sending the an initial delta.

A big drawback is that the synchronization doesn't happen in real-time. This doesn't seem to fit well into the overall Roadster architecture, which is completely event-driven (no polls or "sleeps").

In case the supernode is a HA pair, this variant would generate duplicated traffic. To avoid this, another pair of sockets has to be introduced to synchronize persistence between a HA pair. This also means designing another protocol, and more moving parts overall.

Overall, this variant is very simple, but doesn't offer some features we'd normally expect from a framework like Roadster. The fruits are hanging low; achieving real-time synchronization is easy.

## **PUSH-PULL**

This variant avoids the delays introduced by the polling mechanism of the first variant.

Procedure (for each subnode):

1. via a ROUTER/DEALER socket pair:
  - (a) supernode tells subnode its most recent timestamp in an ICANHAZ request
  - (b) subnode sends delta
  - (c) supernode receives and processes the complete delta
2. subnode sends updates to supernode via PUSH-PULL
3. during low-traffic times, we can send HUGZ as heartbeats

This seems nice at first, but PUSH socket's send buffer will fill up when the connection is interrupted. This isn't bad in and of itself, because when it's full (and writes start to block), we can just destroy the socket and reinitialize and start syncing anew (from ICANHAZ) after a certain timeout. But the problem is that, in case the delta is large, it will inevitably fill the PUSH socket's send buffer, temporarily reaching its high water mark, which is part of its normal operation.

So we'd have to introduce logic to recognize whether the PUSH socket is just temporarily full (e.g. during delta transmission), or permanently full (e.g. the supernode or the link to it is down).

Another disadvantage is that there needs to be another channel to synchronize persistence updates to the other HA peer, if there is one. This means another pair of sockets, another protocol to be designed, and more moving parts overall.

## **PUB/SUB**

This is similar to CSP/CHP. It's not 100% reliable, but even with unstable links, no data loss will occur if the client (the supernode) is able to reconnect within a specific amount of time. ØMQ's default for that amount is 10 seconds. As the requirements specify, 100% consistency is not mandatory for the persistent data.

A possible drawback is that the traffic is duplicated in case the supernode is a HA pair. However, there are numerous opportunities to mitigate this.

Procedure (for each subnode):

1. supernode subscribes to updates from subnode
2. via a ROUTER/DEALER socket pair:
  - (a) supernode tells subnode its most recent timestamp in an ICANHAZ request

- (b) subnode sends delta
  - (c) supernode receives and processes the complete delta
3. supernode starts reading updates, possibly skipping the first few (based on timestamp)

## Chosen Variant

We'll most likely go with the PUB-SUB variant, since it's simple and is similar to what's used for the new CSP in conjunction with multi-node HA. It provides the best opportunities to improve efficiency later on.

Its possible performance issues can be ignored right now, as trying to fix them is arguably considered premature optimization. If this turns out to be an issue in a productive deployment, like over a cellular network link, a future version can switch to multicast. ØMQ supports PGM, which is a reliable multicast protocol. (Pragmatic General Multicast, standardized, directly on top of IP, requires access to raw sockets and thus may require additional privileges) and EPGM (Encapsulated Pragmatic General Multicast, encapsulated in a series of UDP datagrams, doesn't require additional privileges, useful in a ØMQ-only setup).

If its reliability turn out to be an issue, one the socket option `ZMQ_RECOVERY_IVL` can be increased from 10 seconds to, say, 60 seconds, which gives an unstable link more time to recover before any data loss happens.

TODO: describe reasonable default setting, in case we change ZMQ's default.

## Security

TODO briefly describe ØMQ's security features, what's left for us to decide (key distribution)

TODO how it can be verified (-j using wireshark)

TODO why curvezmq is awesome (confidentiality, integrity, availability (server keeps no state until client is fully authenticated))

## OPC UA Interface: High Availability

TODO This is the optional goal.

TODO explain new opportunity for OPC UA HA server

TODO describe whatever needs to be described

- study standard
- use Andy's gem
- according to Andy, this should be a simple thing

## Chapter 4

# Results

TODO what are the results (without discussing them)  
TODO these is probably the "Implementation"

### Port

TODO explain results here

### Cluster

TODO explain results here

### High Availability

TODO explain results here

### Single Level HA

TODO explain results here

### Multi Level HA

TODO explain results here

## **Persistence Synchronization**

TODO explain results here

## **Security**

TODO explain results here

## **OPC UA Interface: High Availability**

TODO explain results here



# Chapter 5

## Discussion

TODO something like a SWOT analysis here (strengths, weaknesses, opportunities, threats)  
TODO general advice: be concise, brief, and specific

## Value Added

TODO what's better than before

## Limitations

TODO identify potential limitations and weaknesses of the product

BStar pair has to be complete (both nodes running) during initialization. Otherwise, only primary node can serve requests; the backup node can't.

Message traffic towards root node sums up because of persistence synchronization. This shouldn't be a problem because of ØMQ's brilliant message batching, so the real limit is given by the inter-node network links.

## Business Benefits

TODO potential applications (UeLS powered by Roadster?)

## Ideas for Improvement

- HA within a node: kill and respawn an actor when it's unresponsive
- switch to Moneta for a unified key-value store interface, then eventually away from Tokyo-Cabinet to something more modern and maintained, like LMDB (it's super fast and crash-proof)

- TIPC: high performance cluster communication protocol, suitable because Roadster nodes are Linux and there are direct links to peers (required for TIPC)
- client authentication
- key management in a DB (instead of files), with GUI to accept new clients
- dynamic node topology (maybe via DSL-file in Etcd, or DIM-only, or Zookeeper)
- other method for data serialization (like MessagePack), would allow adding other programming languages to the cluster
- fast compression for messages, like LZ4 or Snappy
- SERVER/CLIENT sockets from ZMQ 4.2 for simplified message routing

## Chapter 6

# Conclusion

TODO write conclusion, we're the best and everything is awesome TODO they should teach the actor model in APF, because ...

# Bibliography

- [1] B. FIXME Batsov. *Ruby Style Guide*. URL: <https://github.com/bbatsov/ruby-style-guide>.
- [2] A. Dworak, F. Ehm, P. Charrue, and W. Sliwinski. „The new CERN Controls Middleware“. In: *Journal of Physics: Conference Series* 396.012017 (2012). URL: <http://iopscience.iop.org/article/10.1088/1742-6596/396/1/012017/pdf>.

# Part II

## Appendix

## Appendix A

# Self Reflection

TODO how did we perform, completion of goals, accuracy of estimated efforts, efficiency, resourcefulness

## Appendix B

# Task Description

TODO here goes the printed, signed, and scanned Task Description

# Appendix C

## License

As stated in the task description, all of our code contributions underlie the ISC license, which is functionally equivalent to the MIT license and the Simplified BSD license, but uses simpler language. In addition to that, we hereby explicitly grant mindclue GmbH unrestricted usage of all our code contributions.



## Appendix D

# Project Plan

TODO import project plan from wiki  
TODO import risks from wiki

## Organization

TODO roles, how we organize ourselves and how we communicate with each other

# Appendix E

## ZMQ

TODO explain ZMQ in greater detail

TODO strong abstraction (one socket for many connections, connection handling transparent, transport and encryption transparent, no concept of peer addresses)

TODO brokerless/with broker, up to you

TODO basic patterns

TODO extended patterns

TODO not only a "MOM", but a multi threading library (Actor pattern)

## Appendix F

# Infrastructural Problems

TODO describe serious problems here, if any

## Project Management Software

TODO Github/Trello/Harvest/Everhour/Elegantt/Ganttify/Redmine