

BACHELOR THESIS

Roadster High Availability

Manuel Schuler, Patrik Wenger

for industry client
mindclue GmbH

supervised by
Prof. Farhad Mehta

Fall semester 2016

Abstract

TODO introduction

TODO approach and technologies

TODO result

Declaration of Originality

We hereby confirm that we are the sole authors of this document and the described changes to the Roadster framework and libraries developed.

TODO any usage agreements or license

Acknowledgements

TODO anyone we'd like to thank

Management Summary

Initial Situation

TODO describe initial situation, not too technical

Roadster is a next generation monitoring application.

Software Development Process

TODO describe decision to use RUP/Scrum

TODO maybe describe what project management tools we'll be using

Personal Goals

TODO describe personal goal: the cztop-patterns gem

Project Phases

TODO describe this phase in retrospection

Inception

TODO include Gantt chart for this phase

TODO describe this phase in retrospection

Elaboration

TODO include Gantt chart for this phase

TODO describe this phase in retrospection

Construction

TODO include Gantt chart for this phase

TODO describe this phase in retrospection

Transition

TODO include Gantt chart for this phase

TODO describe this phase in retrospection

Results

TODO describe results

Contents

I	Technical Report	1
1	Scope	2
	Motivation	2
	Initial Situation	2
	ZeroMQ	2
	Software Architecture	3
	Goals	3
2	Requirements	4
	Priorities	4
	Functional	4
	Cluster	4
	Single Level HA	5
	Multi Level HA	5
	Persistence Synchronization	5
	Security	5
	OPC UA HA	6
	Use Cases	6
	Non-Functional Requirements	6
3	Methodology	7
	Port to new ZMQ library	7
	Cluster	7
	High Availability	8
	Persistence Synchronization	8
	Security	8
	OPC UA Interface: High Availability	9
4	Results	10
	Port	10
	Cluster	10
	High Availability	10
	Persistence Synchronization	11
	Security	11
	OPC UA Interface: High Availability	11
5	Discussion	12
6	Conclusion	13
II	Appendix	14
A	Self Reflection	15

B Task Description	16
C Project Plan	17
Organization	17
D ZMQ	18
E Infrastructural Problems	19
Project Management Software	19

List of Figures

List of Tables

Listings

Part I

Technical Report

Chapter 1

Scope

TODO what's this thesis about

Motivation

TODO Why do we care about this thesis? Why are we interested?

Initial Situation

TODO What's Roadster and its goals

ZeroMQ

To be able to understand Roadster's architecture, it's helpful to understand the basics of ZeroMQ (hereafter ZMQ) first. This is a brief introduction to ZMQ for the unfamiliar reader.

ZMQ is a message oriented middleware implemented as a library (it doesn't require a broker). Instead, it provides a set of unique socket types and a very abstract, simple interface to interact with them. Socket types:

REQ and REP strict request-reply message handling (synchronous and thus fragile and un-flexible)

DEALER and ROUTER similar, but doesn't enforce strict request-reply fashion (asynchronous and thus more robust and flexible)

PUB and SUB publish-subscribe

PUSH and PULL fan in/fan out, useful for pipelines, round robbin on sender side, fair queueing on receiver side

TODO not strictly server/client type: order of bind/connect doesn't matter

TODO binding/connecting a socket to multiple endpoints is possible

TODO a single socket for many connections, connection handling is completely abstracted away

TODO supported transports: in-memory, IPC, TCP, multicast (UDP or TCP?), TIPC

TODO to achieve reliability or any custom message flow desirable, you build your own solution using this library, see zguide.

TODO doesn't provide serialization (use protobuf/msgpack/json/...)

ZMQ is very light-weight and incorporates advanced technologies such as intelligent message batching and zero-copying which allow it to offer very high throughputs and extremely low latencies. It is several orders of magnitude more performant than other well known MOMs, including closed source and open source solutions.

TODO Roadster architecture



TODO mandatory goals

TODO optional goals

Chapter 2

Requirements

TODO the requirements

Priorities

In descending priority:

1. multi-node CSP
2. single-level HA
3. multi-level HA
4. persistence synchronization
5. security
6. OPC UA HA (optional)

The following sections explain the requirements in greater detail.

Functional

Cluster

This could also be called "Multi-node CSP".

- this is to allow running Roadster in a hierarchical setup
- new COMM actors for inter node communication
- usually 2 (or 3) levels of Roadster nodes
- common cases:
 - - single level, single node (legacy)
 - - single level HA
 - - multi level, HA at root node only
- exotic cases:
 - - multi level, HA at bottom
 - - multi level, HA in middle

- every subtree can live on autonomously
- only node A has write access to values on A (to avoid uncertain situations involving race conditions), e.g.:
 - - a forced value coming from the web UI comes through a command,
 - - routed to the relevant node, where it is applied,
 - - and then synced (up via DEALER and down via PUB, we suppose)
- KISS



Single Level HA

This is where there's a node pair directly connected to a PLC. Both nodes have read/write access to the PLC, but only one of the nodes (the active one) must do so. The nodes must automatically find consensus on who's active. The passive one must automatically take over in case the active one is confirmed to be dead.

TODO the kinds of failures we want to be able to handle: exactly hardware/software failure of the primary node, and network failure (stated by the Task Description)

Multi Level HA

This is where a node pair is the parent of one or more other nodes (subnodes).

TODO the kinds of failures we want to be able to handle: exactly hardware/software failure of the primary node, and network failure (stated by the Task Description)

Persistence Synchronization

This is about the synchronization of the TokyoCabinet databases. Data flow is from south to north (towards the root node), so the root node collects and maintains a replication of the persisted data of all subnodes, recursively.

- autonomous
- not same as CHP
- data only flows from bottom to top

Security

- transport needs to be secure (encrypted and authenticated)
- TODO verify requirements with Andy (we didn't really discuss this during the meeting)

- this requirement comes as the last mandatory goal not because it's insignificant, but because it's easy to enable transport level security on ZMQ sockets, and it would just interfere with the previous development

OPC UA HA

- provide standardized interface upwards from HA pair

Use Cases

TODO maybe there are any?

Non-Functional Requirements

TODO the NFRs

Testing

- we write unit tests for our own contributions
- we test the integrated result in a close-to-reality setup

Coding Guidelines

- basically Ruby style guide¹
- method calls: only use parenthesis when needed, even with arguments (as opposed to ²)
- 2 blank lines before method definition (slightly extending ³)
- YARD API doc, 1 blank comment line before param documentation, one blank comment line before code (ignoring ⁴)
- Ruby 1.9 symbol keys are wanted (just like ⁵)
- align multiple assignments so there's a column of equal signs

¹<https://github.com/bbatsov/ruby-style-guide>

²<https://github.com/bbatsov/ruby-style-guide#method-invocation-parens>

³<https://github.com/bbatsov/ruby-style-guide#empty-lines-between-methods>

⁴<https://github.com/bbatsov/ruby-style-guide#rdoc-conventions>

⁵<https://github.com/bbatsov/ruby-style-guide#hash-literals>

Chapter 3

Methodology

TODO what have we done to arrive at the goal (should be reproducible)
TODO this is probably what we know as "Concept"

Port to new ZMQ library

TODO justify why port is needed right at the beginning (exclude faults from unmaintained ffi-rmq gem, encryption is needed anyway, all the other tasks involve communication over ZMQ)
TODO explain binding options out there, why CZTop (including difference between ZMQ and CZMQ)
TODO explain preliminary task of adding support for the ZMQ options FD and EVENTS in CZTop
TODO explain concept of exchanging ffi-rmq with CZTop

Cluster

TODO explain planned multi node setup
TODO election/design of appropriate protocol
TODO explain Clustered Hashmap Protocol (I guess)

- PCP: use DIM to know node tree and determine next hop for (dialog or fire+forget) messages
- decide on sync variant
 - variant 1
 - * always sync on self-subtree only
 - * con: no copy of remaining tree
 - variant 2:
 - * always sync on complete tree
 - * get snapshot and merge own subtree
 - variant 3:

- * make it configurable: either sync on subtree or complete tree
- node topology in DSL, static file shared on all nodes, read by each actor on startup
- specific config file on each node (conf.rb) knows its own place in topology


```
conf.system_id = "nodes.root"#no HA
OR
conf.system_id = "nodes.root_ha.foo"#with root HA, subnode A directly below root level
```
- maybe a HA pair is one DIM object, has one name, but two IP addresses (primary and backup, in order)

High Availability

TODO we have two different kinds of HA

TODO explain how the failures we're required to be able to handle can be handled

TODO explain similarities between the two kinds of HA

Single Level

- this is different from what's described in the zguide because the concept of client requests is missing here (PLCs don't request anything)
- life sign from one node to the other through some continually updated PLC value
- mark active HA peer in DIM, OR PUSH-PULL & different route back
- side note: PUSH-PULL is probably not feasible, because message are sent to inactive pull anyway, until queue full

Multi Level

TODO explain why is this one different from SL-HA

TODO Finding consensus should be easier here, as it's closely related to the CHP described in the zguide.

Persistence Synchronization

- super node requests for delta of TC periodically

Security

TODO briefly describe ZMQ's security features, what's left for us to decide (key distribution)

TODO how it can be verified (-i using wireshark)

OPC UA Interface: High Availability

TODO This is the optional goal.

TODO explain new opportunity for OPC UA HA server

TODO describe whatever needs to be described

- study standard
- use Andy's gem
- according to Andy, this should be a simple thing

Chapter 4

Results

TODO what are the results (without discussing them)
TODO these is probably the "Implementation"

Port

TODO explain results here

Cluster

TODO explain results here

High Availability

TODO explain results here

Single Level HA

TODO explain results here

Multi Level HA

TODO explain results here

Persistence Synchronization

TODO explain results here

Security

TODO explain results here

OPC UA Interface: High Availability

TODO explain results here

Chapter 5

Discussion

TODO identify potential limitations and weaknesses of the product
TODO potential applications (UeLS on Roadster?)
TODO be concise, brief, and specific

Chapter 6

Conclusion

TODO write conclusion, we're the best and everything is awesome

Part II

Appendix

Appendix A

Self Reflection

TODO how did we perform, completion of goals, accuracy of estimated efforts, efficiency, resourcefulness

Appendix B

Task Description

TODO here goes the printed, signed, and scanned Task Description

Appendix C

Project Plan

TODO import from wiki

Organization

TODO roles, how we organize ourselves and how we communicate with each other

Appendix D

ZMQ

TODO explain ZMQ in greater detail

TODO strong abstraction (one socket for many connections, connection handling transparent, transport and encryption transparent, no concept of peer addresses)

TODO brokerless/with broker, up to you

TODO basic patterns

TODO extended patterns

TODO not only a "MOM", but a multi threading library (Actor pattern)

Appendix E

Infrastructural Problems

TODO describe serious problems here, if any

Project Management Software

TODO Github/Trello/Harvest/Everhour/Elegantt/Ganttify/Redmine